

Lab Exercise 8 – Terraform Multiple tfvars Files

Objective:

Learn how to use multiple tfvars files in Terraform for different environments.

Prerequisites:

- Terraform installed on your machine.
- Basic knowledge of Terraform configuration and variables.

Steps:

1. Create a Terraform Directory:

```
mkdir terraform-multiple-tfvars  
cd terraform-multiple-tfvars
```

- Create Terraform Configuration Files:
- Create a file named main.tf:

```
# main.tf
```

```
provider "aws" {  
    region = var.region  
}  
  
resource "aws_instance" "example" {  
    ami      = var.ami  
    instance_type = var.instance_type  
}
```

- Create a file named variables.tf:

variables.tf

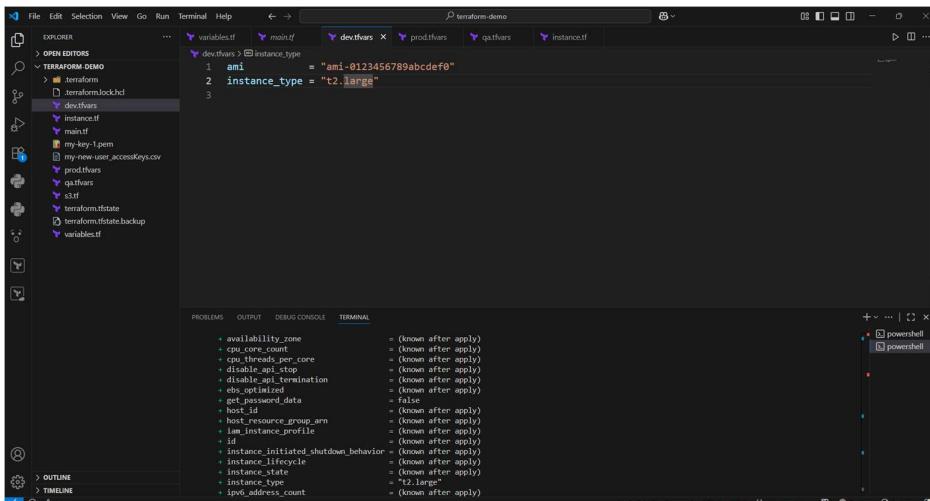
```
variable "ami" {
  type = string
}

variable "instance_ty" {
  type = string
}
```

2. Create Multiple tfvars Files:

- Create a file named dev.tfvars:

dev.tfvars



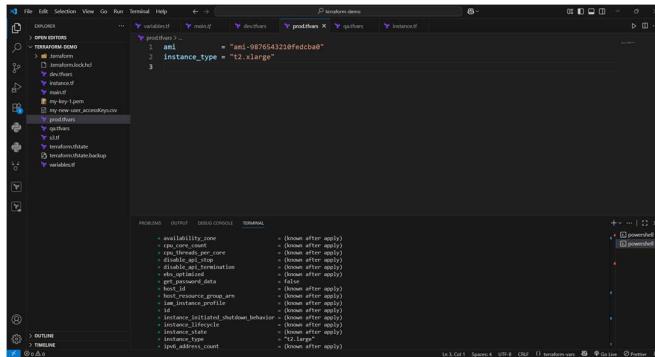
The screenshot shows the VS Code interface with the 'devtfvars' file open in the center editor tab. The code contains two variable assignments:

```
1 ami      = "ami-0123456789abcdef0"
2 instance_type = "t2.large"
```

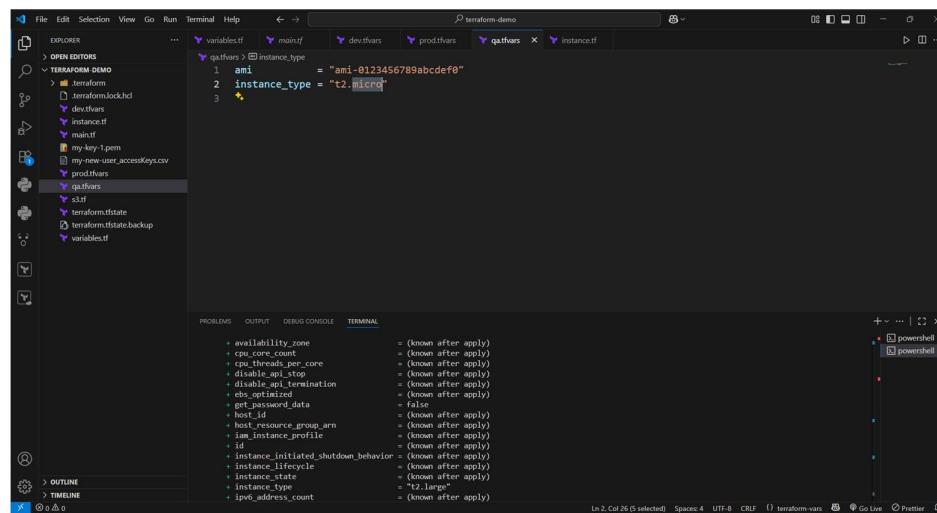
The left sidebar shows the project structure with files like 'variables.tf', 'main.tf', 'devtfvars', 'instance.tf', 'main.tf', 'my-new-user-accessKeys.csv', 'qatfvars', 's3tf', and 'terraform.tfstate'. The bottom status bar indicates 'Create a'.

- file named prod.tfvars:

prod.tfvars



- In these files, provide values for the variables based on the environments.



3. Initialize and Apply for Dev Environment:

- Run the following Terraform commands to initialize and apply the configuration for the dev environment:

```
terraform init
```

```
terraform apply -var-file=dev.tfvars
```

4. Initialize and Apply for Prod Environment:

- Run the following Terraform commands to initialize and apply the configuration for the prod environment:

```
# P:\ C:\Users\hp\Desktop\Notes\DevSecOps-Lab>terraform plan -var-file="dev.tfvars"
# +x
aws_s3_bucket.my_bucket: Refreshing state... [id=s3-bucket-1-terraform]
Terraform used the selected providers to generate the following execution plan. Resource actions are indicated with the following symbols:
+ create
Terraform will perform the following actions:

# aws_instance.web2 will be created
resource "aws_instance" "web2" {
  ami                               = "ami-0123456789abcdef0"
  arn                             = (known after apply)
  associate_public_ip_address      = (known after apply)
  availability_zone                = (known after apply)
  cpu_core_count                   = (known after apply)
  cpu_threads_per_core            = (known after apply)
  ebs_optimized                    = (known after apply)
  disable_api_termination         = (known after apply)
  ebs_optimized_data              = (known after apply)
  host_id                          = (known after apply)
  host_resource_group_arn          = (known after apply)
  iam_instance_profile             = (known after apply)
  id                                = (known after apply)
  instance_initiated_shutdown_behavior = (known after apply)
  instance_type                    = (known after apply)
  instance_type_recycle           = (known after apply)
  instance_state                  = (known after apply)
  instance_type                = "t2.large"
  ipv4_address_count              = (known after apply)
  ipv6_address_count              = (known after apply)
  key_name                         = (known after apply)
  monitoring                       = (known after apply)
  capacity_reservationSpecification = (known after apply)
  password_data                   = (known after apply)
  placement_group                 = (known after apply)
  placement_group_id              = (known after apply)
  primary_network_interface_id    = (known after apply)
  private_dns                      = (known after apply)
  private_ip                       = (known after apply)
  public_ip                        = (known after apply)
  public_ip                         = (known after apply)
  secondary_private_ip             = (known after apply)
  security_group_ids              = (known after apply)
  source_dest_check               = true
  spot_instance_request_id        = (known after apply)
  subnet_id                        = (known after apply)
  tags {
    Name = "Ec2-Instance-Terraform"
  }
  tags_all = {
    Name = "Ec2-Instance-Terraform"
  }
  tenancy                         = (known after apply)
  user_data                        = (known after apply)
  user_data_base64                = (known after apply)
  user_data_replace_on_change     = false
  vpc_security_group_ids          = (known after apply)
  capacity_reservation_specification = (known after apply)
  cpu_options                      = (known after apply)
  ebs_block_device                 = (known after apply)
  enclave_options                 = (known after apply)
  ephemeral_block_device           = (known after apply)
  instance_market_options          = (known after apply)
  maintenance_options              = (known after apply)
  metadata_options                 = (known after apply)
  network_interface                = (known after apply)
  private_dns_name_options         = (known after apply)
  root_block_device                = (known after apply)
}

Plan: 1 to add, 0 to change, 0 to destroy.

Note: You didn't use the -out option to save this plan, so Terraform can't guarantee to take exactly these actions if you run "terraform apply" now.
```

```

PS C:\Users\hp\Desktop\Notes\DeviceOps-Lab\terraform-demo> terraform plan -var-file="qa.tfvars"
aws_s3_bucket.my_bucket: Refreshing state... [id=devs-bucket-t-terraform]
Terraform used the selected providers to generate the following execution plan. Resource actions are indicated with the following symbols:
  + create
  - destroy
  ~ update

Terraform will perform the following actions:

  * resource "aws_instance" "web" {
      + ami                                = "ami-0f114607bb996d9f"
      + arn                                = (known after apply)
      + associate_public_ip_address        = (known after apply)
      + availability_zone                  = (known after apply)
      + cpu_core_count                     = (known after apply)
      + cpu_threads_per_core              = (known after apply)
      + display_name                      = (known after apply)
      + disable_amazon_enhanced_monitoring_log = (known after apply)
      + ephemeral_block_device            = (known after apply)
      + host_id                            = (known after apply)
      + host_resource_group                = (known after apply)
      + iam_instance_profile               = (known after apply)
      + instance_initiated_shutdown_behavior = (known after apply)
      + instance_lifecycle                = (known after apply)
      + instance_type                     = "t2.micro"
      + ipv4_address                      = (known after apply)
      + ipv6_address                      = (known after apply)
      + key_name                           = (known after apply)
      + monitoring                         = (known after apply)
      + output_arn                         = (known after apply)
      + placement_group                   = (known after apply)
      + placement_partition_number        = (known after apply)
      + primary_network_interface_id     = (known after apply)
      + private_dns                        = (known after apply)
      + private_ip                         = (known after apply)
      + public_ip                          = (known after apply)
      + security_groups                   = (known after apply)
      + source_ip                          = (known after apply)
      + spot_instance_request_id          = (known after apply)
      + subnet_id                          = (known after apply)
      + tags {
          + "Name" = "EC2-Instance-Terraform"
        }
      + tags_all {
          + "Name" = "EC2-Instance-Terraform"
        }
    }
  - tenancy
  - user_data
  - user_data_base64
  - user_data_replace_on_change = false
  - vpc_security_group_ids

  * capacity_reservation_specification (known after apply)

  * ebs_block_device (known after apply)
  * enclave_options (known after apply)
  * ephemeral_block_device (known after apply)
  * instance_market_options (known after apply)
  * maintenance_options (known after apply)
  * metadata_options (known after apply)
  * network_interface (known after apply)
  * private_dns_name_options (known after apply)
  * root_block_device (known after apply)

}
Plan: 1 to add, 0 to change, 0 to destroy.

```

```

PS C:\Users\hp\Desktop\Notes\DeviceOps-Lab\terraform-demo> terraform plan -var-file="prod.tfvars"
aws_s3_bucket.my_bucket: Refreshing state... [id=devs-bucket-t-terraform]
Terraform used the selected providers to generate the following execution plan. Resource actions are indicated with the following symbols:
  + create
  - destroy
  ~ update

Terraform will perform the following actions:

  * resource "aws_instance" "web" {
      + ami                                = "ami-0f114607bb996d9f"
      + arn                                = (known after apply)
      + associate_public_ip_address        = (known after apply)
      + availability_zone                  = (known after apply)
      + cpu_core_count                     = (known after apply)
      + cpu_threads_per_core              = (known after apply)
      + display_name                      = (known after apply)
      + disable_amazon_enhanced_monitoring_log = (known after apply)
      + ephemeral_block_device            = (known after apply)
      + host_id                            = (known after apply)
      + host_resource_group                = (known after apply)
      + iam_instance_profile               = (known after apply)
      + instance_initiated_shutdown_behavior = (known after apply)
      + instance_lifecycle                = (known after apply)
      + instance_type                     = "t2.micro"
      + ipv4_address                      = (known after apply)
      + ipv6_address                      = (known after apply)
      + key_name                           = (known after apply)
      + monitoring                         = (known after apply)
      + output_arn                         = (known after apply)
      + placement_group                   = (known after apply)
      + placement_partition_number        = (known after apply)
      + primary_network_interface_id     = (known after apply)
      + private_dns                        = (known after apply)
      + private_ip                         = (known after apply)
      + public_ip                          = (known after apply)
      + security_groups                   = (known after apply)
      + source_ip                          = (known after apply)
      + spot_instance_request_id          = (known after apply)
      + subnet_id                          = (known after apply)
      + tags {
          + "Name" = "EC2-Instance-Terraform"
        }
      + tags_all {
          + "Name" = "EC2-Instance-Terraform"
        }
    }
  - tenancy
  - user_data
  - user_data_base64
  - user_data_replace_on_change = false
  - vpc_security_group_ids

  * capacity_reservation_specification (known after apply)

  * ebs_block_device (known after apply)
  * enclave_options (known after apply)
  * instance_market_options (known after apply)
  * maintenance_options (known after apply)
  * metadata_options (known after apply)
  * network_interface (known after apply)

}

```

5. Test and Verify:

- Observe how different tfvars files are used to set variable values for different environments during the apply process.

- Access the AWS Management Console or use the AWS CLI to verify the creation of resources in the specified regions and instance types.

6. Clean Up:

- After testing, you can clean up resources:

```
terraform destroy -var-file=dev.tfvars  
terraform destroy -var-file=prod.tfvars
```

- Confirm the destruction by typing yes.

7. Conclusion:

This lab exercise demonstrates how to use multiple tfvars files in Terraform to manage variable values for different environments. It allows you to maintain separate configuration files for different environments, making it easier to manage and maintain your infrastructure code. Experiment with different values in the dev.tfvars and prod.tfvars files to observe how they impact the infrastructure provisioning process for each environment.